Cloud Computing Final Project Demo

Team 08

Model description (Fine tuned XLNet)

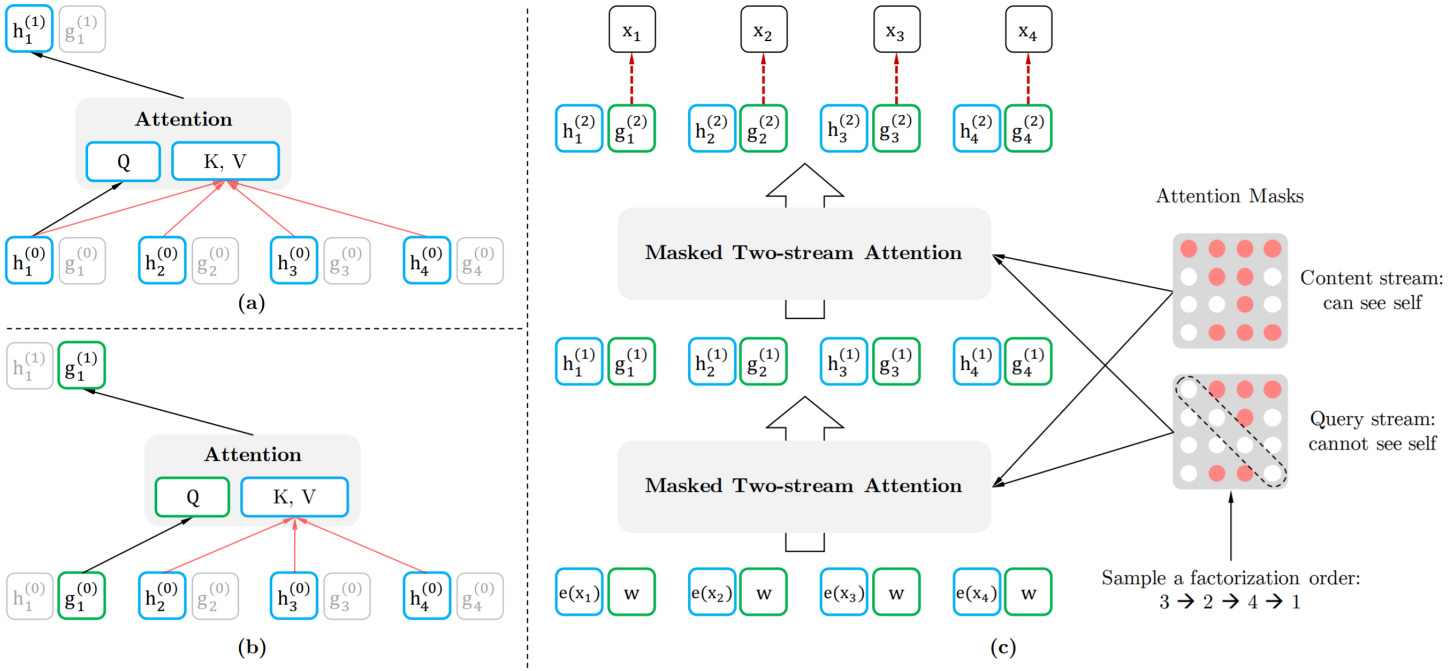

Figure 2: (a): Content stream attention, which is the same as the standard self-attention. (b): Query stream attention, which does not have access information about the content x_{z_t}. (c): Overview of the permutation language modeling training with two-stream attention.

System Structure (Ideal)

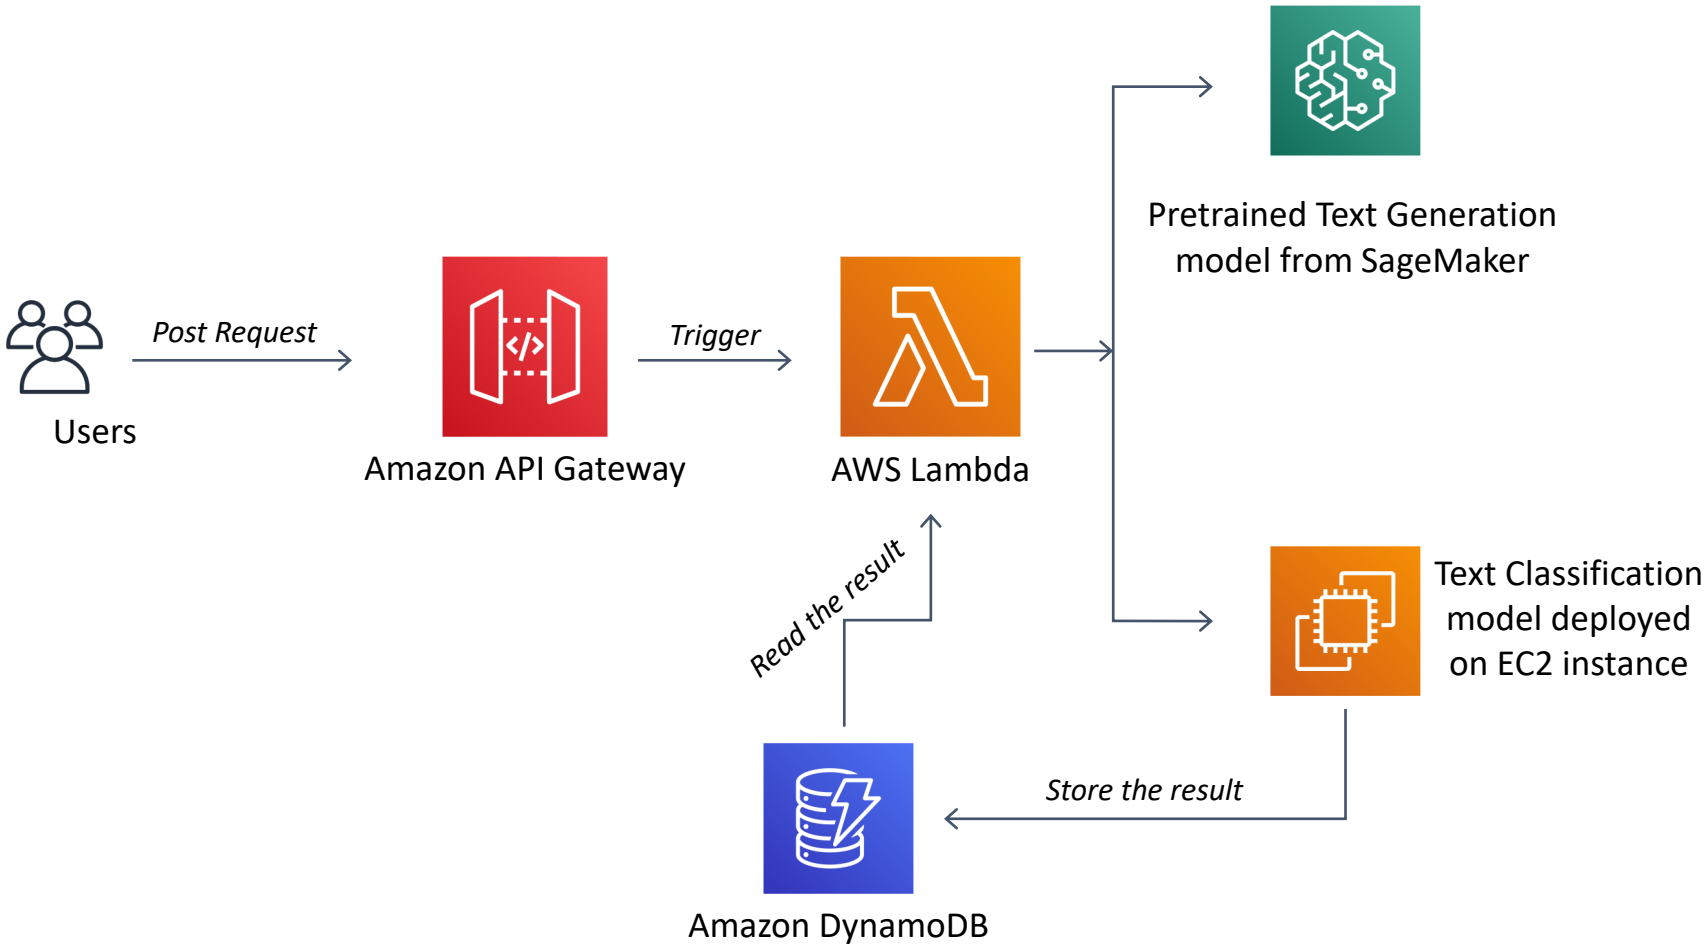

id	result
29146	0
162417	0
226080	[1]
245468	[0]
258456	1
497517	[0]
529481	1
609092	1
628218	[0]
750054	1
792284	1

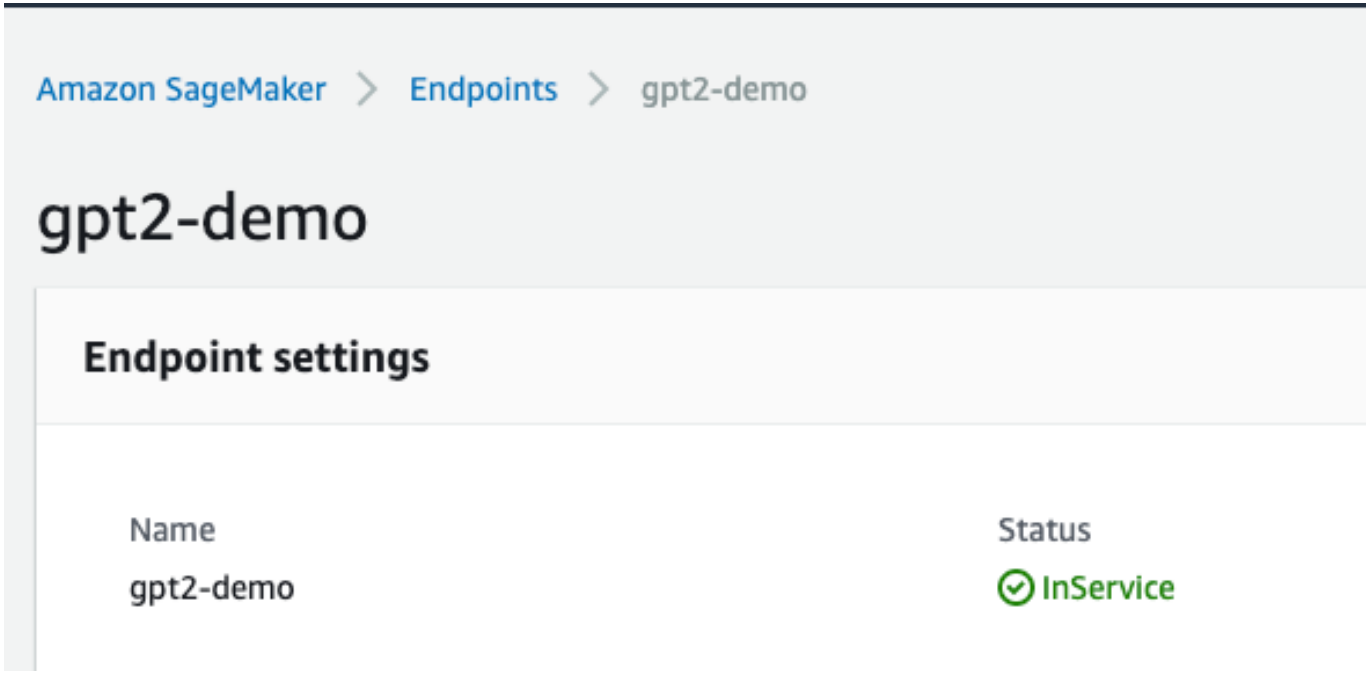
Amazon SageMaker
Endpoints
gpt2-demo
gpt2-demo
Endpoint settings
Name
gpt2-demo
Status
InService

Q1: What's the purpose of the database?

- EC2 instance is not always available.
- The Lambda would first need to **turn on** the instance when requests arrive, and then sends a shell script to the instance (through SSM client).
- The script would
  - Activate the proper environment
  - Run the classification script
- The script will store the result in DynamoDB with a key from the Lambda.

Q2: Why not perform classification inside the Lambda

- XLNet is a HUGE model, which cannot easily fit into Lambda's limited space.
- Deploying `PyTorch`, `Transformers` on Lambda has some tricky dependency problems.
- Lambda does not have GPU.

Practical Issue

- For some reasons, we cannot create IAM role to make components interact with each other.
- So we replace the Lambda and the API Gateway with another EC2 instance, which runs a Flask server.
- Meanwhile, SSM agent is replaced by another Flask server hosted on inference server.

Fake Lambda =_=

```
(base) ubuntu@ip-172-31-23-131:~/lambda$ tree
.
|-- classfication_callback.py
|-- __pycache__
|   |-- classfication_callback.cpython-37.pyc
|   `-- sagemaker_callback.cpython-37.pyc
|-- sagemaker_callback.py
`-- server.py

1 directory, 5 files
```

Future work

- Deploy custom classification task on SageMaker
- Accelerate/Scale up Giant Language Model inference.
- Taking real-world issues into account, e.g. high concurrency.